


Hurricane Katrina Study

by Asad Imam and Michael Burrell



Abstract

In this study, we looked at several of the most devastating hurricanes to hit the United States, such as Hurricane Katrina. We set out to predict what factors (features) associated with hurricanes lead to the most damage (USD) using machine learning models including linear regression, decision trees/random forests, support vector machines, and artificial neural networks. Our results included selection of the Random Forest Tree Ensemble model with an normalized mean residual squared error (NMRSE) of 0.084.

Introduction

Although planet Earth is a beautiful home, we are constantly reminded of the annual occurrence of natural disasters, Hurricanes being one of the world's most recurrent. Each hurricane season (June - November), the United States is hit with about ten storms; 6 of which become hurricanes and about 2 of which become major hurricanes (Category 3 or greater). These hurricanes are responsible for the highest number of deaths (6,593) from 1980-2020 and they cause enormous amounts of damage. Of the 258 U.S. weather disasters since 1980, tropical cyclones have caused the most damage at \$945 billion total, with an average cost of almost \$21.5 billion per event.

The main objective of our project is to analyze historical Hurricane data from numerous sources in order to understand why some hurricanes, such as Katrina in 2005, are more devastating than others by identifying which variables are best for predicting hurricane damage using machine

learning. A major component to why Hurricane Katrina was so devastating is because of the extremely high storm surges that were reported, which led to the breaking of the levees and floodwalls. Storm Surge is measured as feet above the normal tide and is a common statistic used for Hurricane reporting. It is our hypothesis that the higher the storm surge of a hurricane, the higher the damage costs will be. We will explore this problem using a methodology following the 4 Types of Data Analytics; Descriptive Analysis, Diagnostic Analysis, Predictive Analysis, Prescriptive Analysis.

Understanding the what, why, when, where, and how of your data helps drive analytics maturity.



DESCRIPTIVE

What have we done in the past?



DIAGNOSTIC

Why have we seen past results?



PREDICTIVE

Where are we going and when?



PRESCRIPTIVE

How should we take action?

The use of descriptive analytics will help identify certain trends, characteristics, and relationships in the data. We then will use diagnostic analytics to discover why we have seen past results using predictive models on our hypothesized predictor (storm surge) From this, we can then use the same predictive modeling techniques to predict the damage that current and future hurricanes will have based on the features selected. Furthermore, we aim to use prescriptive analytics to help to prepare for such severe impacts, as Katrina caused. This includes faster severe hurricane detection and building smarter infrastructure for the future.

About the Data Set

For our data set we initially considered 30 of the deadliest hurricanes in US History. This data was based on the list published by CBS News article [Deadliest hurricanes in U.S. history](https://www.cbsnews.com/pictures/deadliest-hurricanes-worst-in-the-us-list/). (<https://www.cbsnews.com/pictures/deadliest-hurricanes-worst-in-the-us-list/>). We then added 20 more hurricanes which were also Category 3 or above from the National Hurricane Center's (NHC) National Oceanic and Atmospheric Administration (NOAA) Website. [Link to Website](https://www.nhc.noaa.gov/data/tcr/index.php?season=2005&basin=atl) (<https://www.nhc.noaa.gov/data/tcr/index.php?season=2005&basin=atl>). We collected the following data about the hurricanes

1. **Name** : Name of the Hurricane
2. **Minimum Elevation** : Refers to the Minimum Elevation of the place worst impacted by the Hurricane. This is relative to the sea level (0ft).
3. **Category**: The Category of the Hurricane. (Ranges from 1-5: The hurricanes we considered were all atleast Category 3).
4. **Year** : Year when the Hurricane occurred.
5. **Month** : Month of the Hurricane. (Majority of the Hurricane when hurricane lasted more than a month)
6. **Damage (2021 USD)**: The Damage caused by the Hurricane. This amount is in 2021 USD.
7. **Casualties**: The number of people that died from the Hurricane.
8. **Days Long**: The number of days the Hurricane Lasted.
9. **Minimum Pressure**: The minimum pressure of the hurricane. This is computed in mb.

10. **Peak Wind Speed:** The Peak Wind Speed of the Hurricane. This is computed in mph.

11. **Max Storm Surge:** The Storm Surge of the Hurricane. This is measured in ft.

Note: Storm Surge is measured as the rise in water level above the normal tidal level.

Data Preprocessing

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
```

```
In [2]: 1 df = pd.read_excel('NHC Historical Storm Data.xlsx')
```

```
In [3]: 1 df.head()
```

Out[3]:

	ID	Name	Minimum Elevation (ft)	Category	Year	Month	Damage (2021 USD)	Casualties	Days Long
0	NaN	Great Galveston	7	4	1900	September	1.156021e+09	8000	20
1	NaN	Maria	4	5	2017	September	1.007176e+11	3057	17
2	NaN	Okeechobee	4	5	1928	September	1.604152e+09	2500	16
3	AL122005	Katrina	-8	5	2005	August	2.247291e+11	1833	9
4	NaN	Cheni�re Caminada	-8	4	1893	October	1.523944e+08	2000	9

Finding out how many rows and columns our dataset has

```
In [4]: 1 df.shape
```

Out[4]: (50, 12)

Dropping the ID Column since it doesn't provide any information about the hurricanes

```
In [5]: 1 df = df.drop('ID', 1)
```

```
In [6]: 1 df.columns
```

```
Out[6]: Index(['Name', 'Minimum Elevation (ft)', 'Category', 'Year', 'Month',
              'Damage (2021 USD)', 'Casualties', 'Days Long',
              'Minimum Pressure (mb)', 'Peak Wind Speed (mph)',
              'Max Storm Surge (ft)'],
              dtype='object')
```

We don't have ID in the dataset now

Removing missing values from the data frame

```
In [7]: 1 df.isna().sum()
```

```
Out[7]: Name 0
        Minimum Elevation (ft) 0
        Category 0
        Year 0
        Month 0
        Damage (2021 USD) 1
        Casualties 0
        Days Long 0
        Minimum Pressure (mb) 0
        Peak Wind Speed (mph) 0
        Max Storm Surge (ft) 0
        dtype: int64
```

We have 1 missing value in the Damage (2021 USD) column. Removing that Record

```
In [8]: 1 df.dropna(axis=0,inplace = True)
```

```
In [9]: 1 ##Looking at the shape of our dataset now
        2 df.shape
```

```
Out[9]: (49, 11)
```

We removed the 1 missing record from our dataset

Normalizing the Damage (2021 USD)

```
In [10]: 1 #Looking at the damage
        2 df['Damage (2021 USD)'].min()
```

```
Out[10]: 5836382.98
```

The smallest Damage value is 5.8 million. We will normalize the damage column by dividing each record by 1,000,000

```
In [11]: 1 df['Damage (2021 USD)']=df['Damage (2021 USD)']/1000000
```

```
In [12]: 1 print("the minimum normalized damage is", df['Damage (2021 USD)'].min()
        the minimum normalized damage is 5.836382980000001 and the maximum norma
        lized damage is 224729.134665
```

Looking at the 5 smallest records for damage

```
In [13]: 1 df.nsmallest(5, ['Damage (2021 USD)'])
```

Out[13]:

	Name	Minimum Elevation (ft)	Category	Year	Month	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)	Peak Wind Speed (mph)
21	1886 Indianola	95	2	1886	August	5.836383	150	19	925	150
5	Sea Islands	0	3	1893	August	30.478889	2000	18	954	185
30	Atlantic- Gulf	0	4	1999	September	36.223409	750	12	927	149
18	Georgia	0	4	1898	September	49.574096	179	11	938	130
19	1875 Indianola	95	3	1875	September	99.749091	800	10	955	115

We see that 1886 Indianola is an outlier for damage with 5 million dollars in damage. We will remove this record

```
In [14]: 1 df = df.drop(21)
```

Looking at the 5 smallest records for damage again

```
In [15]: 1 df.nsmallest(5, ['Damage (2021 USD)'])
```

Out[15]:

	Name	Minimum Elevation (ft)	Category	Year	Month	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)	Peak Wind Speed (mph)
5	Sea Islands	0	3	1893	August	30.478889	2000	18	954	185
30	Atlantic- Gulf	0	4	1999	September	36.223409	750	12	927	149
18	Georgia	0	4	1898	September	49.574096	179	11	938	130
19	1875 Indianola	95	3	1875	September	99.749091	800	10	955	115
20	1906 Florida Keys	2	3	1906	October	126.030206	240	15	953	120

The outlier Damage has now been removed from the dataset

Normalizing the Dataset

Normalizing Storm Surge

```
In [16]: 1 df['Max Storm Surge (ft)'] = (df['Max Storm Surge (ft)']-df['Max Storm
```

Normalizing Wind Speed

```
In [17]: 1 df['Peak Wind Speed (mph)'] = (df['Peak Wind Speed (mph)']-df['Peak Win
```

Normalizing Damage

```
In [18]: 1 df['Damage (2021 USD)'] = (df['Damage (2021 USD)']-df['Damage (2021
```

Normalizing Minimum Elevation

```
In [19]: 1 df['Minimum Elevation (ft)'] = (df['Minimum Elevation (ft)']-df['Minimu
```

Normalizing Pressure

```
In [20]: 1 df['Minimum Pressure (mb)'] = (df['Minimum Pressure (mb)']-df['Minimum
```

Normalizing Casualties

```
In [21]: 1 df['Casualties'] = (df['Casualties']-df['Casualties'].mean())/df['Casua
```

Viewing the complete normalized dataset

In [22]:

```
1 df.head(48)
```

Out[22]:

	Name	Minimum Elevation (ft)	Category	Year	Month	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)	
0	Great Galveston	0.272917	4	1900	September	-0.391140	2.269092	20	-0.142631	-0.2
1	Maria	0.134048	5	2017	September	2.279546	0.635993	17	-1.132973	1.2
2	Okeechobee	0.134048	5	1928	September	-0.379120	0.451968	16	-0.228747	0.6
3	Katrina	-0.421430	5	2005	August	5.606090	0.231601	9	-1.391323	1.2
4	Cheni�re Caminada	-0.421430	4	1893	October	-0.418062	0.286775	9	0.589361	-0.4
5	Sea Islands	-0.051112	3	1893	August	-0.421333	0.286775	18	0.847711	1.6
6	San Ciriaco	-0.051112	4	1899	August	-0.404419	0.899641	31	-0.185689	0.2
7	Audrey	-0.421430	3	1957	June	-0.383657	-0.236556	5	0.503244	-0.8
8	Great Labor Day	0.041468	5	1935	September	-0.368440	-0.239200	12	-1.821906	1.6
10	Great Miami	-0.051112	4	1926	September	-0.380578	-0.251093	12	-0.185689	0.2
11	Grand Isle	-0.051112	3	1909	September	-0.413256	-0.258362	10	0.761594	-1.0
12	1919 Florida Keys	0.041468	4	1919	September	-0.412793	-0.118939	13	-0.314864	0.2
13	New Orleans	-0.421430	4	1915	September	-0.412679	-0.283141	10	-0.142631	0.0
14	Galveston	0.272917	4	1915	August	-0.385723	-0.283141	19	0.244894	-0.6
15	Camille	-0.051112	5	1969	August	-0.172182	-0.289418	8	-1.477439	1.2
16	New England	-0.051112	5	1938	September	-0.262461	-0.289418	14	0.244894	0.6
17	Diane	-0.051112	2	1955	August	-0.193798	-0.313206	16	1.493586	-1.6
18	Georgia	-0.051112	4	1898	September	-0.420820	-0.314858	11	0.158778	-0.6
19	1875 Indianola	4.346419	3	1875	September	-0.419474	-0.109688	10	0.890769	-1.2
20	1906 Florida Keys	0.041468	3	1906	October	-0.418769	-0.294704	15	0.804653	-1.0
22	Mississippi	-0.051112	3	1906	September	-0.406435	-0.329725	11	0.804653	-1.0
23	Cedar Keys	-0.051112	3	1896	September	-0.413741	-0.307259	8	1.106061	-0.8
24	Agnes	-0.051112	1	1972	June	-0.052479	-0.331708	12	1.838053	-2.4
25	Hazel	2.217088	4	1954	October	-0.317658	-0.178739	13	0.158778	-0.4
26	Betsy	-0.051112	4	1965	September	-0.090446	-0.347236	16	0.331011	-0.3
27	Great Atlantic	-0.051112	4	1944	September	-0.380342	-0.352852	7	-0.056514	0.0
28	Carol	-0.051112	3	1954	August	-0.295775	-0.350209	6	0.890769	-1.2

	Name	Minimum Elevation (ft)	Category	Year	Month	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)	
29	Floyd	-0.051112	4	1999	September	-0.135065	-0.345253	12	-0.573214	0.4
30	Atlantic-Gulf	-0.051112	4	1999	September	-0.421178	-0.126207	12	-0.314864	0.1
31	Donna	-0.051112	4	1960	August	-0.178533	-0.253737	15	-0.185689	-0.0
32	Alicia	-4.680091	3	1983	August	-0.200517	-0.367059	6	1.192178	-1.2
33	Gilbert	-0.051112	5	1988	September	-0.236795	-0.268934	11	-1.994139	1.5
34	Hugo 1989	-0.051112	5	1989	September	0.230595	-0.353844	14	-0.702389	0.7
35	Andrew 1992	-0.051112	5	1992	August	1.009635	-0.352522	12	-0.530156	1.2
36	Opal 1995	-0.051112	4	1995	October	-0.195223	-0.353183	10	-0.788506	-0.0
37	Mitch 1998	0.550656	5	1998	October	-0.147684	6.010715	17	-1.262148	1.3
38	Keith 2000	-0.051112	4	2000	September	-0.408519	-0.351531	9	0.201836	-0.3
39	Iris 2001	-0.051112	4	2001	October	-0.411763	-0.362103	5	0.589361	-0.0
40	Isabel 2003	-0.051112	5	2004	September	-0.281919	-0.357147	14	-0.831564	0.9
41	Charley 2004	-0.051112	4	2004	August	0.212784	-0.369041	7	1.622761	0.2
42	Frances 2004	-0.051112	4	2004	August	0.045286	-0.357808	18	0.029603	0.0
43	Ivan 2004	-0.051112	5	2004	September	0.376386	-0.333360	24	0.374069	0.8
44	Jeanne 2004	-0.051112	3	2005	September	-0.391221	0.628725	17	0.675478	-1.0
45	Dennis 2005	-0.051112	4	2005	July	-0.338132	-0.344923	15	-0.185689	-0.9
46	Rita 2005	-0.051112	5	2005	September	-0.026547	-0.334351	9	-1.692731	1.2
47	Wilma 2005	-0.051112	5	2005	October	0.421804	-0.345253	13	-2.252489	1.5
48	Hurricane Ike 2008	-0.051112	4	2008	September	-0.319622	-0.309572	15	2.225578	-1.4
49	Ida	-0.051112	4	2021	August	2.126175	-0.336003	10	-0.185689	0.2

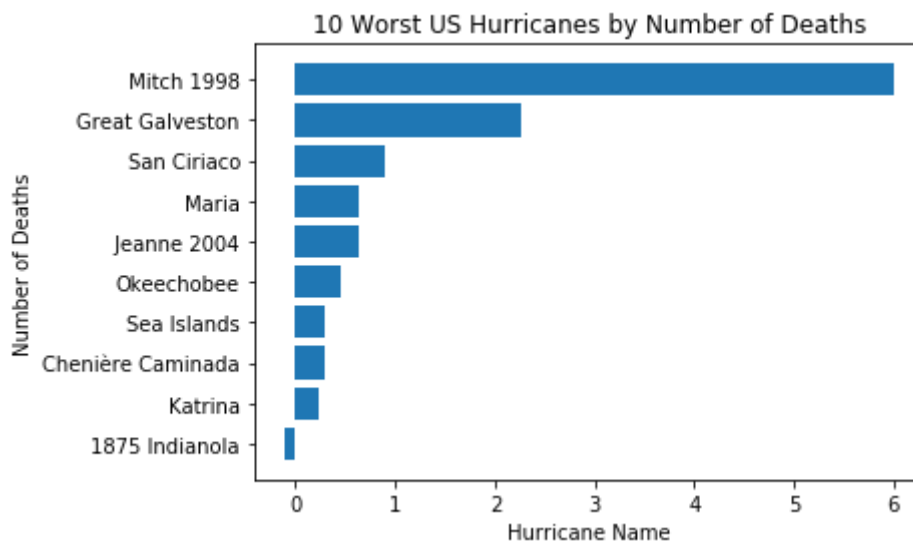
Descriptive Analytics

Descriptive analytics is the interpretation of historical data to better understand changes that have occurred in a business. Descriptive analytics describes the use of a range of historic data to draw comparisons.

Worst US Hurricanes by Number of Deaths

```
In [23]: 1 worst_casualties = df.sort_values('Casualties',ascending=False)[:10]
        2 worst_casualties.sort_values('Casualties',ascending=True, inplace = True)
```


```
In [24]: 1 plt.barh(worst_casualties['Name'],worst_casualties['Casualties'])
        2 plt.title('10 Worst US Hurricanes by Number of Deaths')
        3 plt.xlabel('Hurricane Name')
        4 plt.ylabel('Number of Deaths')
        5 plt.show()
```



Longest US Hurricanes by Number of Days

```
In [25]: 1 longest_hurricanes = df.sort_values('Days Long',ascending=False)[:10]
2 longest_hurricanes.sort_values('Days Long',ascending=True, inplace = True)
```

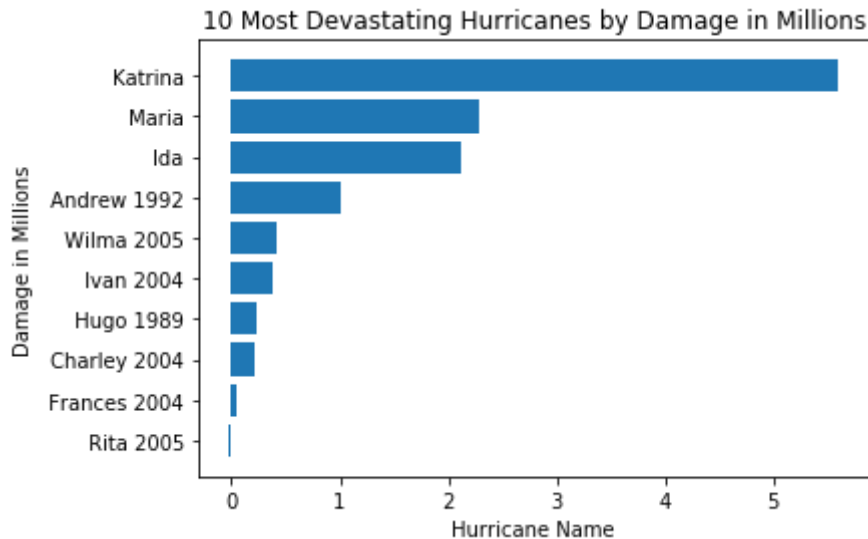
```
In [26]: 1 plt.barh(longest_hurricanes['Name'],longest_hurricanes['Days Long'])
2 plt.title('10 Longest US Hurricanes by Number of Days')
3 plt.xlabel('Hurricane Name')
4 plt.ylabel('Number of Days')
5 plt.show()
```



US Hurricanes by Damage in US Dollars 2021

```
In [27]: 1 expense_hurricanes = df.sort_values('Damage (2021 USD)',ascending=False)
2 expense_hurricanes.sort_values('Damage (2021 USD)',ascending=True, inplace = True)
```

```
In [28]: 1 plt.barh(expense_hurricanes['Name'],expense_hurricanes['Damage (2021 U
2 plt.title('10 Most Devastating Hurricanes by Damage in Millions')
3 plt.xlabel('Hurricane Name')
4 plt.ylabel('Damage in Millions')
5 plt.show()
```

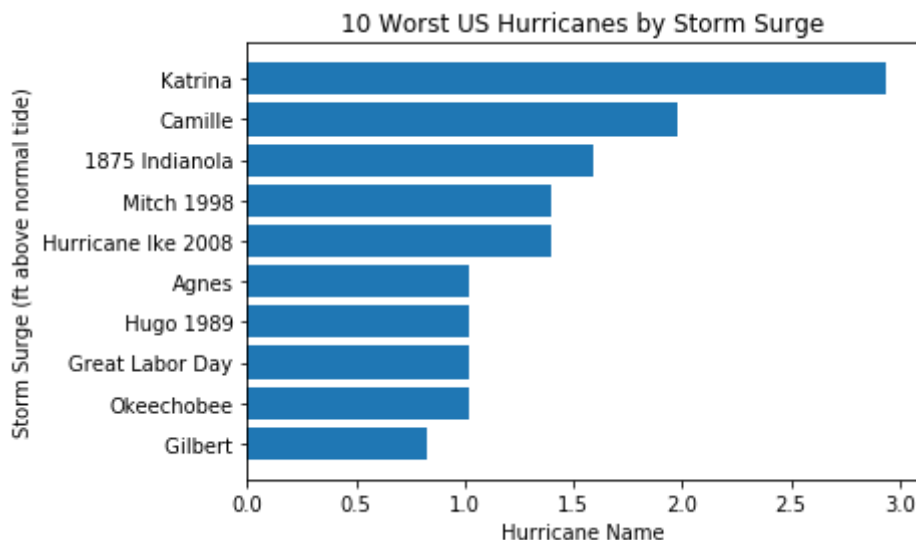


Here we can see that Hurricane Katrina was the most devastating hurricane in terms of damage

Worst US Hurricanes by Storm Surge

```
In [29]: 1 surge_hurricanes = df.sort_values('Max Storm Surge (ft)',ascending=False)
2 surge_hurricanes.sort_values('Max Storm Surge (ft)',ascending=True, inplace=True)
```

```
In [30]: 1 plt.barh(surge_hurricanes['Name'],surge_hurricanes['Max Storm Surge (ft
2 plt.title('10 Worst US Hurricanes by Storm Surge')
3 plt.xlabel('Hurricane Name')
4 plt.ylabel('Storm Surge (ft above normal tide)')
5 plt.show()
```

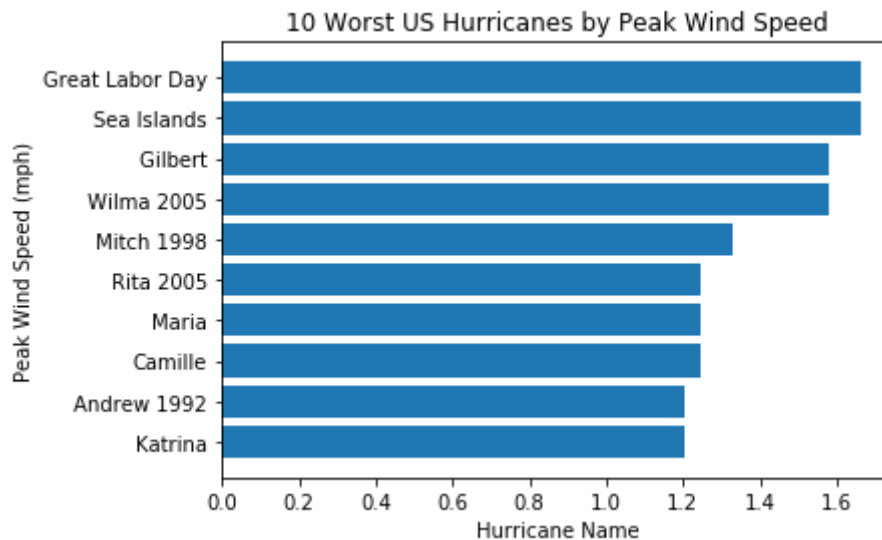


Here we can see that Hurricane Katrina had the highest storm surge

Worst US Hurricanes by Windspeed

```
In [31]: 1 wind_hurricanes = df.sort_values('Peak Wind Speed (mph)',ascending=False)
2 wind_hurricanes.sort_values('Peak Wind Speed (mph)',ascending=True, inplace=True)
```

```
In [32]: 1 plt.barh(wind_hurricanes['Name'],wind_hurricanes['Peak Wind Speed (mph)'])
2 plt.title('10 Worst US Hurricanes by Peak Wind Speed')
3 plt.xlabel('Hurricane Name')
4 plt.ylabel('Peak Wind Speed (mph)')
5 plt.show()
```



Descriptive Statistics

```
In [33]: 1 df.describe()
```

Out[33]:

	Minimum Elevation (ft)	Category	Year	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)
count	4.800000e+01	48.000000	48.000000	4.800000e+01	4.800000e+01	48.000000	4.800000e+01
mean	3.642919e-17	3.979167	1960.666667	9.251859e-18	-1.734723e-17	12.770833	-9.830100e-18
std	1.000000e+00	0.887012	43.471846	1.000000e+00	1.000000e+00	4.938945	1.000000e+01
min	-4.680091e+00	1.000000	1875.000000	-4.213325e-01	-3.690413e-01	5.000000	-2.252489e+01
25%	-5.111165e-02	3.750000	1918.000000	-4.069563e-01	-3.457491e-01	9.750000	-5.409204e-01
50%	-5.111165e-02	4.000000	1967.000000	-3.186398e-01	-3.009818e-01	12.000000	-1.345573e-01
75%	-5.111165e-02	5.000000	2001.750000	-1.239101e-01	-1.656060e-01	15.250000	6.970069e-01
max	4.346419e+00	5.000000	2021.000000	5.606090e+00	6.010715e+00	31.000000	2.225578e+01

Moments

Moments [of a statistical distribution] The shape of any distribution can be described by its various 'moments'. The first four are:

- 1) The mean, which indicates the central tendency of a distribution.
- 2) The second moment is the variance, which indicates the width or deviation.
- 3) The third moment is the skewness, which indicates any asymmetric 'leaning' to either left or right.
- 4) The fourth moment is the Kurtosis, which indicates the degree of central 'peakedness' or, equivalently, the 'fatness' of the outer tails.

```
In [34]: 1 df.mean()
```

```
Out[34]: Minimum Elevation (ft)      3.642919e-17
          Category                    3.979167e+00
          Year                        1.960667e+03
          Damage (2021  USD)          9.251859e-18
          Casualties                  -1.734723e-17
          Days Long                   1.277083e+01
          Minimum Pressure (mb)       -9.830100e-18
          Peak Wind Speed (mph)       -3.978299e-16
          Max Storm Surge (ft)        -1.671117e-16
          dtype: float64
```

```
In [35]: 1 df.var()
```

```
Out[35]: Minimum Elevation (ft)      1.000000
          Category                    0.786791
          Year                        1889.801418
          Damage (2021  USD)          1.000000
          Casualties                  1.000000
          Days Long                   24.393174
          Minimum Pressure (mb)       1.000000
          Peak Wind Speed (mph)       1.000000
          Max Storm Surge (ft)        1.000000
          dtype: float64
```

```
In [36]: 1 df.skew()
```

```
Out[36]: Minimum Elevation (ft)     -0.212971
          Category                   -0.912646
          Year                       -0.371926
          Damage (2021  USD)         4.377213
          Casualties                  5.084860
          Days Long                   1.152828
          Minimum Pressure (mb)      -0.202676
          Peak Wind Speed (mph)      -0.110173
          Max Storm Surge (ft)        0.389234
          dtype: float64
```

```
In [37]: 1 df.kurtosis()
```

```
Out[37]: Minimum Elevation (ft)    17.614028  
         Category                  1.426798  
         Year                     -1.349028  
         Damage (2021 USD)        21.866154  
         Casualties               29.029357  
         Days Long                 2.865794  
         Minimum Pressure (mb)    -0.066244  
         Peak Wind Speed (mph)    -0.566672  
         Max Storm Surge (ft)     0.572863  
         dtype: float64
```

Diagnostic Analytics

Diagnostic analytics is a form of advanced analytics that examines data or content to answer the question, “Why did it happen?” It is characterized by techniques such as drill-down, data discovery, data mining and correlations.

Correlation Matrix

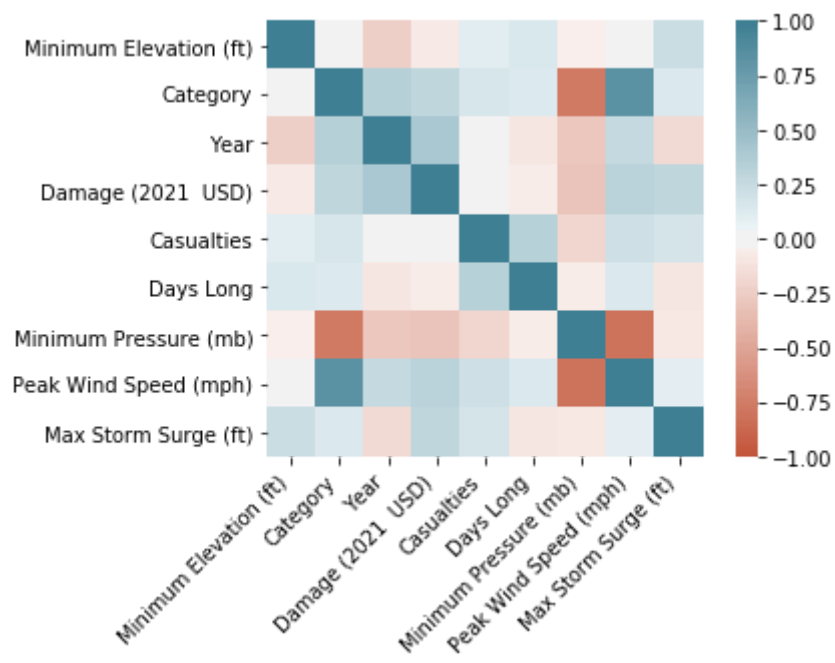
In [38]:

```
1 corr = df.corr()  
2 corr
```

Out[38]:

	Minimum Elevation (ft)	Category	Year	Damage (2021 USD)	Casualties	Days Long	Minimum Pressure (mb)	Peak Wind Speed (mph)	
Minimum Elevation (ft)	1.000000	0.030095	-0.238460	-0.071556	0.110923	0.149988	-0.046375	0.004795	
Category	0.030095	1.000000	0.346333	0.296432	0.170940	0.139730	-0.759842	0.832920	
Year	-0.238460	0.346333	1.000000	0.400277	-0.015717	-0.100947	-0.283216	0.262652	-
Damage (2021 USD)	-0.071556	0.296432	0.400277	1.000000	0.024506	-0.054897	-0.302486	0.322284	
Casualties	0.110923	0.170940	-0.015717	0.024506	1.000000	0.330484	-0.195756	0.214146	
Days Long	0.149988	0.139730	-0.100947	-0.054897	0.330484	1.000000	-0.052042	0.143251	-
Minimum Pressure (mb)	-0.046375	-0.759842	-0.283216	-0.302486	-0.195756	-0.052042	1.000000	-0.803915	-
Peak Wind Speed (mph)	0.004795	0.832920	0.262652	0.322284	0.214146	0.143251	-0.803915	1.000000	
Max Storm Surge (ft)	0.237121	0.141595	-0.157653	0.291497	0.185826	-0.094577	-0.082176	0.091811	

```
In [39]: 1 ax = sns.heatmap(
2         corr,
3         vmin=-1, vmax=1, center=0,
4         cmap=sns.diverging_palette(20, 220, n=200),
5         square=True
6     )
7 ax.set_xticklabels(
8     ax.get_xticklabels(),
9     rotation=45,
10    horizontalalignment='right'
11 );
```



Result: For Diagnostic Analytics, we will predict the value of Storm Surge based on Features such as Minimum Pressure, Minimum Elevation and Peak Wind Speed

Multiple Linear Regression

Using Minimum Pressure and Wind Speed to predict the Storm Surge

```
In [40]: 1 import statsmodels.api as sm
```

```
In [41]: 1 X = df[["Minimum Pressure (mb)", "Peak Wind Speed (mph)"]]
2 y = df["Max Storm Surge (ft)"]
3
4 # Note the difference in argument order
5 model = sm.OLS(y, X).fit()
6 predictions = model.predict(X) # make the predictions by the model
7
8 # Print out the statistics
9 model.summary()
```

Out[41]: OLS Regression Results

Dep. Variable:	Max Storm Surge (ft)	R-squared (uncentered):	0.009
Model:	OLS	Adj. R-squared (uncentered):	-0.034
Method:	Least Squares	F-statistic:	0.2002
Date:	Mon, 20 Dec 2021	Prob (F-statistic):	0.819
Time:	21:55:55	Log-Likelihood:	-67.396
No. Observations:	48	AIC:	138.8
Df Residuals:	46	BIC:	142.5
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Minimum Pressure (mb)	-0.0237	0.247	-0.096	0.924	-0.521	0.473
Peak Wind Speed (mph)	0.0728	0.247	0.295	0.769	-0.424	0.570

Omnibus:	1.482	Durbin-Watson:	1.357
Prob(Omnibus):	0.477	Jarque-Bera (JB):	0.771
Skew:	0.270	Prob(JB):	0.680
Kurtosis:	3.306	Cond. No.	3.03

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Using Minimum Elevation (ft), Minimum Pressure (mb) and Peak Wind Speed (mph) to predict Surge

```
In [42]: 1 X = df[["Minimum Elevation (ft)", "Minimum Pressure (mb)", "Peak Wind Spe
2 y = df["Max Storm Surge (ft)"]
3
4 # Note the difference in argument order
5 model = sm.OLS(y, X).fit()
6 predictions = model.predict(X) # make the predictions by the model
7
8 # Print out the statistics
9 model.summary()
```

Out[42]: OLS Regression Results

Dep. Variable:	Max Storm Surge (ft)	R-squared (uncentered):	0.064
Model:	OLS	Adj. R-squared (uncentered):	0.002
Method:	Least Squares	F-statistic:	1.033
Date:	Mon, 20 Dec 2021	Prob (F-statistic):	0.387
Time:	21:55:55	Log-Likelihood:	-66.005
No. Observations:	48	AIC:	138.0
Df Residuals:	45	BIC:	143.6
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Minimum Elevation (ft)	0.2369	0.145	1.639	0.108	-0.054	0.528
Minimum Pressure (mb)	0.0048	0.243	0.020	0.984	-0.485	0.494
Peak Wind Speed (mph)	0.0946	0.243	0.389	0.699	-0.394	0.584

Omnibus:	2.551	Durbin-Watson:	1.363
Prob(Omnibus):	0.279	Jarque-Bera (JB):	1.598
Skew:	0.270	Prob(JB):	0.450
Kurtosis:	3.712	Cond. No.	3.04

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Computing Errors

Root Mean Squared Error (**RMSE**) and Mean Absolute Error (**MAE**) are metrics used to evaluate a Regression Model. These metrics tell us how accurate our predictions are and, what is the amount of deviation from the actual values.

Technically, RMSE is the Root of the Mean of the Square of Errors and MAE is the Mean of Absolute value of Errors. Here, errors are the differences between the predicted values (values predicted by our regression model) and the actual values of a variable. They are calculated as follows :

$$\text{RMSE} = \sqrt{\frac{\sum (y_i - y_p)^2}{n}}$$

$$\text{MAE} = \frac{|(y_i - y_p)|}{n}$$

y_i = actual value

y_p = predicted value

n = number of observations/rows

```
In [43]: 1 #RMSE/MAE computation using sklearn library
        2 from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [44]: 1 multiple_regression_rmse = round(np.sqrt(mean_squared_error(y, predictions)),3)
        2 multiple_regression_mae = round(mean_absolute_error(y, predictions),3)
        3 print(f'The root mean square error (RMSE) for our multiple regression model is: {multiple_regression_rmse}')
        4 print(f'The mean absolute error (MAE) for our multiple regression model is: {multiple_regression_mae}')
```

The root mean square error (RMSE) for our multiple regression model is:
0.957

The mean absolute error (MAE) for our multiple regression model is: 0.719

The Normalized Root Mean Square Error (**NRMSE**) the RMSE facilitates the comparison between models with different scales. the normalised RMSE (NRMSE) which relates the RMSE to the observed range of the variable. Thus, the NRMSE can be interpreted as a fraction of the overall range that is typically resolved by the model.

There are multiple ways of Normalizing RMSE. See below:

You can normalize by

- the **mean**: $NRMSE = \frac{RMSE}{g}$ (similar to the CV and applied in *INDperform*)
- the **difference between maximum and minimum**: $NRMSE = \frac{RMSE}{y_{max} - y_{min}}$,
- the **standard deviation**: $NRMSE = \frac{RMSE}{\sigma}$, or
- the **interquartile range**: $NRMSE = \frac{RMSE}{Q1 - Q3}$, i.e. the difference between 25th and 75th percentile, of observations.

We will use the difference between the maximum and minimum to normalize the RMSE

```
In [45]: 1 multiple_regression_nrmse = round(multiple_regression_rmse/(y.max()-y.min()),3)
        2 print(f'The normalized root mean square error (NRMSE) for our multiple regression model is: {multiple_regression_nrmse}')
```

The normalized root mean square error (NRMSE) for our multiple regression model is: 0.199

Support Vector Machine

```
In [46]: 1 #Importing the libraries
        2 from sklearn.svm import SVR
```

Predicting Surge using Minimum Elevation (ft), Minimum Pressure (mb) and Peak Wind Speed (mph) using SVM

```
In [47]: 1 X = df[["Minimum Elevation (ft)", "Minimum Pressure (mb)", "Peak Wind Spe
        2 y = df[["Max Storm Surge (ft)"]]
```

```
In [48]: 1 #Feature Scaling
        2 from sklearn.preprocessing import StandardScaler
        3 sc_X = StandardScaler()
        4 sc_y = StandardScaler()
        5 X = sc_X.fit_transform(X)
        6 y = sc_y.fit_transform(y)
```

```
In [49]: 1 #Fitting the Support Vector Regression Model to the dataset
        2 regressor = SVR()
        3 regressor.fit(X,y)
```

/Users/asadimam270/opt/anaconda3/lib/python3.7/site-packages/sklearn/util/s/validation.py:760: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

```
Out[49]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='scal
e',
        kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
In [50]: 1 predictions = regressor.predict(X)
```

```
In [51]: 1 svm_rmse = round(np.sqrt(mean_squared_error(y, predictions)),3)
        2 svm_mae = round(mean_absolute_error(y, predictions),3)
        3 svm_nrmse = round(svm_rmse/(y.max()-y.min()),3)
        4 print(f'The root mean square error (RMSE) for our model is: {svm_rmse}')
        5 print(f'The mean absolute error (MAE) for our model is: {svm_mae}')
        6 print(f'The normalized root mean square error (NRMSE) for our model is:
```

The root mean square error (RMSE) for our model is: 0.895

The mean absolute error (MAE) for our model is: 0.624

The normalized root mean square error (NRMSE) for our model is: 0.184

Tuning the Parameters

```
In [52]: 1 #Fitting the Support Vector Regression Model to the dataset
2 regressor = SVR(C=1000, cache_size=200, coef0=0.0, degree=3, epsilon=0.
3         kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False
4 regressor.fit(X,y)
5 #Predict
6 predictions = regressor.predict(X)
7 #Computing errors
8 svm_rmse = round(np.sqrt(mean_squared_error(y, predictions)),3)
9 svm_mae = round(mean_absolute_error(y, predictions),3)
10 svm_nrmse = round(svm_rmse/(y.max()-y.min()),3)
11 print(f'The root mean square error (RMSE) for our model is: {svm_rmse}')
12 print(f'The mean absolute error (MAE) for our model is: {svm_mae}')
13 print(f'The normalized root mean square error (NRMSE) for our model is:
```

The root mean square error (RMSE) for our model is: 0.636

The mean absolute error (MAE) for our model is: 0.323

The normalized root mean square error (NRMSE) for our model is: 0.131

/Users/asadimam270/opt/anaconda3/lib/python3.7/site-packages/sklearn/util
s/validation.py:760: DataConversionWarning: A column-vector y was passed
when a 1d array was expected. Please change the shape of y to (n_samples,
,) for example using ravel().
y = column_or_1d(y, warn=True)

Decision Trees

Predicting Surge using Minimum Elevation (ft), Minimum Pressure (mb) and Peak Wind Speed (mph) using Decision Tree Model

```
In [53]: 1 # Import the necessary modules and libraries
2 from sklearn.tree import DecisionTreeRegressor
```

```
In [54]: 1 X = df[["Minimum Elevation (ft)", "Minimum Pressure (mb)", "Peak Wind Spe
2 y = df[["Max Storm Surge (ft)"]]
```

```
In [55]: 1 # Fitting the Decision Tree Regression model
2 dt_regr = DecisionTreeRegressor()
3 dt_regr.fit(X, y)
```

```
Out[55]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

```
In [56]: 1 # Predict
2 y_predict = dt_regr.predict(X)
```

```
In [57]: 1 dt_rmse = round(np.sqrt(mean_squared_error(y, y_predict)),3)
2 dt_mae = round(mean_absolute_error(y, y_predict),3)
3 dt_nrmse = round(dt_rmse/(y.max()-y.min()),3)
4 print(f'The root mean square error (RMSE) for our model is: {dt_rmse}')
5 print(f'The mean absolute error (MAE) for our model is: {dt_mae}')
6 print(f'The normalized root mean square error (NRMSE) for our model is:
```

The root mean square error (RMSE) for our model is: 0.126

The mean absolute error (MAE) for our model is: 0.029

The normalized root mean square error (NRMSE) for our model is: Max Storm
Surge (ft) 0.026

dtype: float64

Artificial Neural Networks

Predicting Surge using Minimum Elevation (ft), Minimum Pressure (mb) and Peak Wind Speed (mph) using ANN

```
In [58]: 1 X = df[["Minimum Elevation (ft)", "Minimum Pressure (mb)", "Peak Wind Spe
2 y = df[["Max Storm Surge (ft)"]]
3 sc_X = StandardScaler()
4 X = sc_X.fit_transform(X)
```

```
In [59]: 1 #Importing the libraries
2 from keras.models import Sequential
3 from keras.layers import Dense
4
5 # Initialising the ANN
6 model = Sequential()
7
8 # Adding the input layer and the first hidden layer
9 model.add(Dense(6, activation = 'relu', input_dim = 3))
10
11 # Adding the output layer
12 model.add(Dense(units = 1))
13
14 model.compile(loss = 'mean_squared_error', metrics = ["accuracy"])
15
16 # Training the model
17 model1 = model.fit(X, y, batch_size = 10, epochs = 100)
18
19 # Testing the Model
20 y_pred = model.predict(X)
```

```
Epoch 1/100
5/5 [=====] - 1s 1ms/step - loss: 1.1948 - accur
acy: 0.0000e+00
Epoch 2/100
5/5 [=====] - 0s 1ms/step - loss: 1.1701 - accur
acy: 0.0000e+00
Epoch 3/100
5/5 [=====] - 0s 1ms/step - loss: 1.1526 - accur
acy: 0.0000e+00
Epoch 4/100
5/5 [=====] - 0s 1ms/step - loss: 1.1372 - accur
acy: 0.0000e+00
Epoch 5/100
5/5 [=====] - 0s 2ms/step - loss: 1.1246 - accur
acy: 0.0000e+00
Epoch 6/100
5/5 [=====] - 0s 1ms/step - loss: 1.1135 - accur
acy: 0.0000e+00
Epoch 7/100
5/5 [=====] - 0s 1ms/step - loss: 1.1010 - accur
acy: 0.0000e+00
```

```
In [60]: 1 ann_rmse = round(np.sqrt(mean_squared_error(y, y_pred)),3)
2 ann_mae = round(mean_absolute_error(y, y_pred),3)
3 ann_nrmse = round(ann_rmse/(y.max()-y.min()),3)
4 print(f'The root mean square error (RMSE) for our ANN model is: {ann_rm
5 print(f'The mean absolute error (MAE) for our ANN model is: {ann_mae}')
6 print(f'The normalized root mean square error (NRMSE) for our ANN model
```

```
The root mean square error (RMSE) for our ANN model is: 0.904
The mean absolute error (MAE) for our ANN model is: 0.697
The normalized root mean square error (NRMSE) for our ANN model is: Max S
torm Surge (ft)      0.188
dtype: float64
```


Model Selection

```
In [61]: 1 print("Errors for SVM")
          2 print(f'RMSE : {svm_rmse}')
          3 print(f'MAE : {svm_mae}')
          4 print(f'NRMSE : {svm_nrmse}')
          5 print('')
          6 print("Errors for Decision Tree")
          7 print(f'RMSE : {dt_rmse}')
          8 print(f'MAE : {dt_mae}')
          9 print(f'NRMSE : {dt_nrmse}')
         10 print('')
         11 print("Errors for ANN")
         12 print(f'RMSE : {ann_rmse}')
         13 print(f'MAE : {ann_mae}')
         14 print(f'NRMSE : {ann_nrmse}')
```

Errors for SVM

RMSE : 0.636

MAE : 0.323

NRMSE : 0.131

Errors for Decision Tree

RMSE : 0.126

MAE : 0.029

NRMSE : Max Storm Surge (ft) 0.026

dtype: float64

Errors for ANN

RMSE : 0.904

MAE : 0.697

NRMSE : Max Storm Surge (ft) 0.188

dtype: float64

```

In [62]: 1 # import module
2 from tabulate import tabulate
3
4 # assign data
5 mydata = [{"Decision Tree", "RMSE", dt_rmse, 0.89},
6           ["Decision Tree", "MAE", dt_mae, 0.67],
7           ["Decision Tree", "RMSE", dt_nrmse, 1.85e-17],
8           ["SVM", "RMSE", svm_rmse, 0.96],
9           ["SVM", "MAE", svm_mae, 0.70],
10          ["SVM", "RMSE", svm_nrmse, 0.11],
11          ["ANN", "RMSE", ann_rmse, 0.96],
12          ["ANN", "MAE", ann_mae, 0.72],
13          ["ANN", "RMSE", ann_nrmse, 0.1]]
14
15 # create header
16 head = ["Model", "Measure", "Error in Python", "Error in R"]
17
18 # display table
19 print('Errors for our models in Python and R:')
20 print(tabulate(mydata, headers=head, tablefmt="grid"))

```

Errors for our models in Python and R:

Model	Measure	Error in Python	Error in R
Decision Tree	RMSE	0.126	0.89
Decision Tree	MAE	0.029	0.67
Decision Tree	RMSE	0.026	1.85e-17
SVM	RMSE	0.636	0.96
SVM	MAE	0.323	0.7
SVM	RMSE	0.131	0.11
ANN	RMSE	0.904	0.96
ANN	MAE	0.697	0.72
ANN	RMSE	0.188	0.1

Result: From the errors above we can see that Decision Tree was our best model to predict the Storm Surge

Predictive Analytics

Predictive analytics encompasses a variety of statistical techniques from data mining, predictive modelling, and machine learning that analyze current and historical facts to make predictions about future or otherwise unknown events.

We will be using the same models as above to predict the Damage based on Storm Surge, Minimum Elevation, Minimum Pressure and Peak Wind Speed

Support Vector Machines

```
In [63]: 1 X = df[["Max Storm Surge (ft)", "Minimum Elevation (ft)", "Minimum Pressu
2 y = df["Damage (2021 USD)"]
```

```
In [64]: 1 #Feature Scaling
2 from sklearn.preprocessing import StandardScaler
3 sc_X = StandardScaler()
4 sc_y = StandardScaler()
5 X = sc_X.fit_transform(X)
6 y = sc_y.fit_transform(y.values.reshape(-1,1))
```

```
In [65]: 1 #Fitting the Support Vector Regression Model to the dataset
2 regressor = SVR()
3 regressor.fit(X,y)
```

```
/Users/asadimam270/opt/anaconda3/lib/python3.7/site-packages/sklearn/util
s/validation.py:760: DataConversionWarning: A column-vector y was passed
when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
Out[65]: SVR(C=1.0, cache_size=200, coef0=0.0, degree=3, epsilon=0.1, gamma='scal
e',
          kernel='rbf', max_iter=-1, shrinking=True, tol=0.001, verbose=False)
```

```
In [66]: 1 predictions = regressor.predict(X)
```

```
In [67]: 1 svm_rmse_pred = round(np.sqrt(mean_squared_error(y, predictions)),3)
2 svm_mae_pred = round(mean_absolute_error(y, predictions),3)
3 svm_nrmse_pred = round(svm_rmse_pred/(y.max()-y.min()),3)
4 print(f'The root mean square error (RMSE) for our model is: {svm_rmse_p
5 print(f'The mean absolute error (MAE) for our model is: {svm_mae_pred}'
6 print(f'The normalized root mean square error (NRMSE) for our model is:
```

The root mean square error (RMSE) for our model is: 0.934

The mean absolute error (MAE) for our model is: 0.341

The normalized root mean square error (NRMSE) for our model is: 0.153

Artificial Neural Networks

```
In [68]: 1 X = df[["Max Storm Surge (ft)", "Minimum Elevation (ft)", "Minimum Pressu
2 y = df["Damage (2021 USD)"]
```

```
In [69]: 1 sc_X = StandardScaler()
2 X = sc_X.fit_transform(X)
```

```
In [70]: 1 #Importing the libraries
2 from keras.models import Sequential
3 from keras.layers import Dense
4
5 # Initialising the ANN
6 model = Sequential()
7
8 # Adding the input layer and the first hidden layer
9 model.add(Dense(8, activation = 'relu', input_dim = 4))
10
11 # Adding the output layer
12 model.add(Dense(units = 1))
13
14 model.compile(loss = 'mean_squared_error', metrics=["accuracy"])
15
16 # Training the model
17 model1 = model.fit(X, y, batch_size = 10, epochs = 100)
18
19 # Testing the Model
20 y_pred = model.predict(X)
```

```
Epoch 1/100
5/5 [=====] - 0s 1ms/step - loss: 1.0580 - accur
acy: 0.0000e+00
Epoch 2/100
5/5 [=====] - 0s 1ms/step - loss: 1.0257 - accur
acy: 0.0000e+00
Epoch 3/100
5/5 [=====] - 0s 1ms/step - loss: 1.0062 - accur
acy: 0.0000e+00
Epoch 4/100
5/5 [=====] - 0s 1ms/step - loss: 0.9895 - accur
acy: 0.0000e+00
Epoch 5/100
5/5 [=====] - 0s 1ms/step - loss: 0.9770 - accur
acy: 0.0000e+00
Epoch 6/100
5/5 [=====] - 0s 1ms/step - loss: 0.9644 - accur
acy: 0.0000e+00
Epoch 7/100
5/5 [=====] - 0s 1ms/step - loss: 0.9520 - accur
acy: 0.0000e+00
```

```
In [71]: 1 ann_rmse_pred = round(np.sqrt(mean_squared_error(y, y_pred)),3)
2 ann_mae_pred = round(mean_absolute_error(y, y_pred),3)
3 ann_nrmse_pred = round(ann_rmse_pred/(y.max()-y.min()),3)
4 print(f'The root mean square error (RMSE) for our ANN model is: {ann_rm
5 print(f'The mean absolute error (MAE) for our ANN model is: {ann_mae_pr
6 print(f'The normalized root mean square error (RMSE) for our ANN model
```

```
The root mean square error (RMSE) for our ANN model is: 0.867
The mean absolute error (MAE) for our ANN model is: 0.459
The normalized root mean square error (RMSE) for our ANN model is: 0.459
```

Decision Tree

```
In [72]: 1 X = df[["Max Storm Surge (ft)", "Minimum Elevation (ft)", "Minimum Pressu
2 y = df["Damage (2021 USD)"]
```

```
In [73]: 1 dt_regr = DecisionTreeRegressor()
2 dt_regr.fit(X, y)
```

```
Out[73]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort='deprecated',
                                random_state=None, splitter='best')
```

```
In [74]: 1 # Predict
2 y_predict = dt_regr.predict(X)
```

```
In [75]: 1 dt_rmse_pred = np.sqrt(mean_squared_error(y, y_predict))
2 dt_mae_pred = mean_absolute_error(y, y_predict)
3 dt_nrmse_pred = dt_rmse_pred/(y.max()-y.min())
4 print(f'The root mean square error (RMSE) for our model is: {dt_rmse_pr
5 print(f'The mean absolute error (MAE) for our model is: {dt_mae_pred}')
6 print(f'The normalized root mean square error (NRMSE) for our model is:
```

```
The root mean square error (RMSE) for our model is: 4.951616820006514e-05
The mean absolute error (MAE) for our model is: 1.0107445508998844e-05
The normalized root mean square error (NRMSE) for our model is: 8.2151480
71977484e-06
```

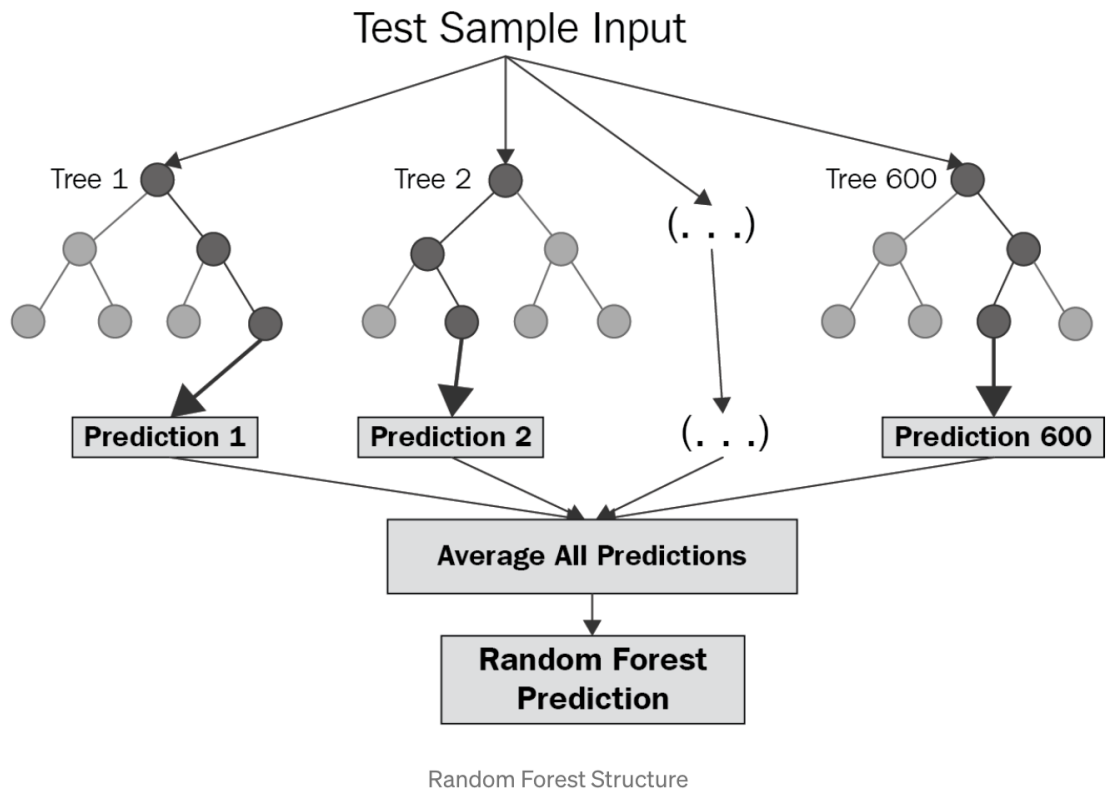
Result: It seems that our Decision Tree Model is overfitting. We will use the Tree Ensemble Method instead.

Theory: "Decision trees carry a big risk of overfitting, and tend to find local optima because they can't go back after they have made a split. To address these weaknesses, we turn to Random Forest, which illustrates the power of combining many decision trees into one model."

This info was taken from the Article below: [Random Forest Regression](https://medium.com/swlh/random-forest-and-its-implementation-71824ced454f)
(<https://medium.com/swlh/random-forest-and-its-implementation-71824ced454f>)

Ensemble Method: Ensemble means a group of elements viewed as a whole rather than individually. An Ensemble method creates multiple models and combines them to solve it. Ensemble methods help to improve the robustness/generalizability of the model.

Random Forrest



```
In [76]: 1 x = df[["Max Storm Surge (ft)", "Minimum Elevation (ft)", "Minimum Pressu
2 y = df["Damage (2021 USD)"]
```

```
In [77]: 1 from sklearn.ensemble import RandomForestRegressor
```

For this model I've chosen 5 trees (n_estimator=5)

```
In [78]: 1 rf_pred = RandomForestRegressor(n_estimators=5)
```

```
In [79]: 1 #Training the Model
2 rf_pred.fit(X, y)
3 # Predict
4 y_predict = rf_pred.predict(X)
```

```
In [80]: 1 rf_rmse_pred = round(np.sqrt(mean_squared_error(y, y_predict)),3)
2 rf_mae_pred = round(mean_absolute_error(y, y_predict),3)
3 rf_nrmse_pred = round(rf_rmse_pred/(y.max()-y.min()),3)
4 print(f'The root mean square error (RMSE) for our model is: {rf_rmse_pr
5 print(f'The mean absolute error (MAE) for our model is: {rf_mae_pred}')
6 print(f'The normalized root mean square error (NRMSE) for our model is:
```

The root mean square error (RMSE) for our model is: 0.496

The mean absolute error (MAE) for our model is: 0.241

The normalized root mean square error (NRMSE) for our model is: 0.082

Model Selection

```
In [81]: 1 # assign data
2 errors_pred = [
3     ["SVM", "RMSE", svm_rmse_pred],
4     ["SVM", "MAE", svm_mae_pred],
5     ["SVM", "NRMSE", svm_nrmse_pred],
6     ["ANN", "RMSE", ann_rmse_pred],
7     ["ANN", "MAE", ann_mae_pred],
8     ["ANN", "NRMSE", ann_nrmse_pred],
9     ["Decision Tree", "RMSE", dt_rmse_pred],
10    ["Decision Tree", "MAE", dt_mae_pred],
11    ["Decision Tree", "NRMSE", dt_nrmse_pred],
12    ["Random Forrest", "RMSE", rf_rmse_pred],
13    ["Random Forrest", "MAE", rf_mae_pred],
14    ["Random Forrest", "NRMSE", rf_nrmse_pred],
15
16 ]
17
18 # create header
19 head = ["Model", "Measure", "Error in Python"]
20
21 # display table
22 print(tabulate(errors_pred, headers=head, tablefmt="grid"))
```

Model	Measure	Error in Python
SVM	RMSE	0.934
SVM	MAE	0.341
SVM	NRMSE	0.153
ANN	RMSE	0.867
ANN	MAE	0.459
ANN	NRMSE	0.144
Decision Tree	RMSE	4.95162e-05
Decision Tree	MAE	1.01074e-05
Decision Tree	NRMSE	8.21515e-06
Random Forrest	RMSE	0.496
Random Forrest	MAE	0.241
Random Forrest	NRMSE	0.082

Results: Although Decision Trees returned the lowest errors, we suspect that this might be an issue of overfitting (which is a common error with Decision Tree). To resolve this we used an ensemble approach (namely Random Forrest). Based on the errors above we will say that Random Forrest was our best model for predicting the Damage.

Conclusion

The main goal of this study was to use machine learning algorithms to identify which variables or features are best for predicting hurricane damage. Through descriptive and diagnostic analytics, we found that Max Storm Surge, Minimum Elevation, Minimum Pressure, and Peak Wind Speed were the best predictors for our target, Damage. Our predictive analytics included modeling using linear regression, decision trees/random forests, support vector machines, and artificial neural networks. The selection of the Random Forest Tree Ensemble model with an normalized mean residual squared error (NMRSE) of 0.084 was our best predictive model, thus the best model to predict damage for any hurricane in history or the future.

Discussion

Limitations to our study included data availability and accuracy issues. Some of our hurricane records went back as far as the 1800s and there just isn't much data on hurricanes that far back. Additionally, we felt that the integrity of our data would be lost if we decided to add more features with inconsistent and/or inaccurate data. We see a great opportunity for future work on this study, especially on the addition of new, accurate, and insightful features. This can include but is not limited to precipitation, sea surface temperature, and tornadoes.

References

1. <https://worldpopulationreview.com/state-rankings/states-with-the-least-natural-disasters> (<https://worldpopulationreview.com/state-rankings/states-with-the-least-natural-disasters>)
 2. <https://coast.noaa.gov/states/fast-facts/hurricane-costs.html#:~:text=Of%20the%2058%20U.S.%20weather,6%2C593%20between%201980%2> (<https://coast.noaa.gov/states/fast-facts/hurricane-costs.html#:~:text=Of%20the%2058%20U.S.%20weather,6%2C593%20between%201980%2>)
 3. <https://www.cnn.com/2021/04/13/weather/2021-atlantic-hurricane-season-fast-facts/index.html> (<https://www.cnn.com/2021/04/13/weather/2021-atlantic-hurricane-season-fast-facts/index.html>)
 4. <https://www.cbsnews.com/pictures/deadliest-hurricanes-worst-in-the-us-list/> (<https://www.cbsnews.com/pictures/deadliest-hurricanes-worst-in-the-us-list/>)
-